

# Geosnap Tutorial

Stephen Francisco, Kate Oxx

April 14, 2021

This tutorial will showcase tools from the Geosnap package for Python. Geosnap helps spatial analysts identify, explore, and model neighborhoods statically or temporally based on parameters chosen by the user. It contains built-in data from the past three censuses (1990, 2000, and 2010) as well a series of other similar datasets containing American geographic and demographic data but also allows the user to import their own data.

In this tutorial we will cover boundary harmonization for standardizing geographic units over time, geodemographics and regionalization, and simulating neighborhood change – just part of the functionality of the Geosnap package.

## 0.0.1 Geosnap Installation & Data Preparation

Installation The makers of Geosnap recommend installing with anaconda. We recommend you try and use the python-unstable env where we installed pysal last week. (Alternately, you might want to use the gus5031 env.) To install, open the Miniconda prompt, navigate to the proper environment, and use the following commands:

```
[ ]: conda activate <your_env>
      conda install -c conda-forge geosnap
```

Be sure to confirm it's installed, or if you installed it previously, make sure you're running version 0.6.0.

## 0.0.2 Built-In Data

Geosnap comes “batteries included” with multiple demographic datasets, including state, county, and micropolitan and metropolitan statistical (MSA) boundaries. In order to view the list of datasets, run the following commands:

```
[3]: from geosnap import datasets
      dir(datasets)
```

```
[3]: ['acs',
      'blocks_2000',
      'blocks_2010',
      'codebook',
      'counties',
      'ltldb',
      'msa_definitions',
      'msas',
      'ncdb',
      'states',
      'tracts_1990',
      'tracts_2000',
      'tracts_2010']
```

Each dataset is a pandas (or geopandas) dataframe, whose attributes and methods can be used.

Next, we will import geosnap’s central data structure: **Community** > “A Community is a dataset that stores information about a collection of neighborhoods over several time periods, including each neighborhood’s physical, socioeconomic, and demographic attributes and its demarcated boundaries. Under the hood, each Community is simply a long-form geopandas geodataframe with some associated metadata.”<sup>1</sup>

```
[4]: from geosnap import Community
```

Processing times can be fairly long when working with large datasets (such as large cities or entire states). Therefore, we recommend storing the census data locally to improve performance using the following commands:

```
[5]: from geosnap.io import store_census
      store_census()
```

Loading manifest:

100% | 8/8

[00:00<00:00, 4.04k/s]

Copying objects:

100% | 253M/253M

[00:07<00:00, 34.3MB/s]

Successfully installed package 'census/tracts\_cartographic', tophash=e848bc9 from s3://spatial-ucr

Loading manifest:

100% | 9/9

[00:00<00:00, 9.03k/s]

Copying objects:

100% | 185M/185M

[00:03<00:00, 51.3MB/s]

Successfully installed package 'census/administrative', tophash=3eecddec from s3://spatial-ucr

To harness the census data built into Geosnap, we will have to assign it to a variable, thus creating our own Community. In this case, we will use the `Community.from_census` constructor to obtain demographic data for our county of interest (Philadelphia) using FIPS county codes:

```
[6]: phl = Community.from_census(county_fips="42101")
```

To access the underlying data, call its `gdf` attribute to return a geodataframe (`gdf`). Here, we use the `head` command to show the first five rows in the `gdf`:

```
[7]: phl.gdf.head()
```

```
[7]:
```

	geoid	n_mexican_pop	n_cuban_pop	n_puerto_rican_pop	\
29740	42101036500	9185.0	9.0	79.0	
29758	42101035800	6012.0	15.0	30.0	
29773	42101035900	5354.0	16.0	39.0	
29775	42101036400	528.0	1.0	17.0	
29785	42101035700	8529.0	15.0	55.0	

	n_total_housing_units	n_vacant_housing_units	\
29740	3684.0	248.0	
29758	2344.0	196.0	
29773	2235.0	82.0	
29775	2.0	0.0	
29785	3869.0	169.0	

	n_occupied_housing_units	n_owner_occupied_housing_units	\
29740	3436.0	2352.0	
29758	2148.0	1865.0	
29773	2153.0	1538.0	
29775	2.0	1.0	
29785	3700.0	1361.0	

	n_renter_occupied_housing_units	n_white_persons	...	\
29740	1084.0	8880102.0	...	
29758	283.0	572268.0	...	
29773	615.0	504466.0	...	
29775	1.0	3233.0	...	
29785	2339.0	796592.0	...	

	p_irish_born_pop	p_italian_born_pop	p_poverty_rate_children	\
29740	NaN	NaN	NaN	
29758	NaN	NaN	NaN	
29773	NaN	NaN	NaN	
29775	NaN	NaN	NaN	
29785	NaN	NaN	NaN	

	p_poverty_rate_hispanic	p_russian_born_pop	p_scandinavian_born_pop	\
29740	NaN	NaN	NaN	
29758	NaN	NaN	NaN	
29773	NaN	NaN	NaN	
29775	NaN	NaN	NaN	
29785	NaN	NaN	NaN	

	p_scandinavian_pop	n_total_pop_sample	p_female_labor_force	\
29740	NaN	NaN	NaN	
29758	NaN	NaN	NaN	
29773	NaN	NaN	NaN	
29775	NaN	NaN	NaN	
29785	NaN	NaN	NaN	

	p_black_persons
29740	NaN
29758	NaN
29773	NaN
29775	NaN
29785	NaN

[5 rows x 195 columns]

In order to view all column names, we can iterate through the series and print each column name:

```
[8]: columns = phl.gdf.columns
      for col in columns:
          print (col)
```

```
geoid
n_mexican_pop
n_cuban_pop
n_puerto_rican_pop
n_total_housing_units
n_vacant_housing_units
n_occupied_housing_units
n_owner_occupied_housing_units
n_renter_occupied_housing_units
n_white_persons
n_nonhisp_white_persons
n_black_persons
n_nonhisp_black_persons
n_hispanic_persons
n_native_persons
n_hawaiian_persons
n_asian_indian_persons
n_chinese_persons
n_filipino_persons
```

n\_japanese\_persons  
n\_korean\_persons  
n\_asian\_persons  
n\_vietnamese\_persons  
n\_white\_age\_distribution  
n\_white\_under\_15  
n\_white\_over\_60  
n\_white\_over\_65  
n\_black\_age\_distribution  
n\_black\_under\_15  
n\_black\_over\_60  
n\_black\_over\_65  
n\_hispanic\_age\_distribution  
n\_hispanic\_under\_15  
n\_hispanic\_over\_60  
n\_hispanic\_over\_65  
n\_native\_age\_distribution  
n\_native\_under\_15  
n\_native\_over\_60  
n\_native\_over\_65  
n\_asian\_age\_distribution  
n\_asian\_under\_15  
n\_total\_pop  
n\_russian\_pop  
n\_italian\_pop  
n\_german\_pop  
n\_irish\_pop  
n\_foreign\_born\_pop  
n\_recent\_immigrant\_pop  
n\_naturalized\_pop  
n\_age\_5\_older  
n\_other\_language  
n\_limited\_english  
median\_home\_value  
median\_contract\_rent  
n\_structures\_30\_old  
n\_occupied\_housing\_units\_sample  
n\_household\_recent\_move  
n\_persons\_under\_18  
n\_persons\_over\_60  
n\_persons\_over\_75  
n\_persons\_over\_15  
n\_persons\_over\_25  
n\_married  
n\_widowed\_divorced  
n\_total\_families  
n\_female\_headed\_families  
n\_female\_over\_16

n\_female\_labor\_force  
n\_labor\_force  
n\_unemployed\_persons  
n\_employed\_over\_16  
n\_employed\_professional  
n\_employed\_manufacturing  
n\_employed\_self\_employed  
n\_civilians\_over\_16  
n\_veterans  
n\_civilians\_16\_64  
n\_disabled  
median\_household\_income  
n\_total\_households  
n\_white\_households  
n\_black\_households  
n\_hispanic\_households  
n\_asian\_households  
per\_capita\_income  
n\_poverty\_determined\_persons  
n\_poverty\_persons  
n\_poverty\_over\_65  
n\_poverty\_determined\_families  
n\_poverty\_determined\_white  
n\_poverty\_white  
n\_poverty\_determined\_black  
n\_poverty\_black  
n\_poverty\_determined\_native  
n\_poverty\_native  
n\_poverty\_determined\_asian  
n\_poverty\_asian  
n\_edu\_college\_greater  
n\_edu\_hs\_less  
p\_mexican\_pop  
p\_cuban\_pop  
p\_puerto\_rican\_pop  
p\_russian\_pop  
p\_italian\_pop  
p\_german\_pop  
p\_irish\_pop  
p\_foreign\_born\_pop  
p\_recent\_immigrant\_pop  
p\_naturalized\_pop  
p\_other\_language  
p\_limited\_english  
n\_housing\_units\_multiunit\_structures\_denom  
n\_total\_housing\_units\_sample  
p\_vacant\_housing\_units  
p\_owner\_occupied\_units

p\_structures\_30\_old  
p\_household\_recent\_move  
p\_persons\_under\_18  
p\_persons\_over\_60  
p\_persons\_over\_75  
p\_married  
p\_widowed\_divorced  
p\_female\_headed\_families  
p\_nonhisp\_white\_persons  
p\_nonhisp\_black\_persons  
p\_hispanic\_persons  
p\_native\_persons  
p\_asian\_persons  
p\_hawaiian\_persons  
p\_asian\_indian\_persons  
p\_chinese\_persons  
p\_filipino\_persons  
p\_japanese\_persons  
p\_korean\_persons  
p\_vietnamese\_persons  
p\_white\_under\_15  
p\_white\_over\_60  
p\_white\_over\_65  
p\_black\_under\_15  
p\_black\_over\_60  
p\_black\_over\_65  
p\_hispanic\_under\_15  
p\_hispanic\_over\_60  
p\_hispanic\_over\_65  
p\_native\_under\_15  
p\_native\_over\_60  
p\_native\_over\_65  
p\_asian\_under\_15  
p\_edu\_hs\_less  
p\_edu\_college\_greater  
p\_unemployment\_rate  
p\_employed\_professional  
p\_employed\_manufacturing  
p\_employed\_self\_employed  
p\_veterans  
p\_disabled  
p\_poverty\_rate  
p\_poverty\_rate\_over\_65  
p\_poverty\_rate\_white  
p\_poverty\_rate\_black  
p\_poverty\_rate\_native  
p\_poverty\_rate\_asian  
geometry

```

year
median_income_asianhh
median_income_blackhh
median_income_hispanichh
median_income_whitehh
n_asian_over_60
n_asian_over_65
n_civilians_over_18
n_german_born_pop
n_housing_units_multiunit_structures
n_irish_born_pop
n_italian_born_pop
n_poverty_determined_hispanic
n_poverty_families_children
n_poverty_hispanic
n_russian_born_pop
n_scandaniavian__born_pop
n_scandaniavian_pop
p_asian_over_60
p_asian_over_65
p_german_born_pop
p_housing_units_multiunit_structures
p_irish_born_pop
p_italian_born_pop
p_poverty_rate_children
p_poverty_rate_hispanic
p_russian_born_pop
p_scandnavian_born_pop
p_scandnavian_pop
n_total_pop_sample
p_female_labor_force
p_black_persons

```

### 0.0.3 Boundary Harmonization with Areal Interpolation

One of the most fascinating features of Geosnap is its ability to harmonize data within geographical boundaries that have changed shape over time. A common example of this is US census boundary changes that can occur from one census period to the next. These boundary changes make it difficult to accurately compare census areas over time, which makes harmonization an incredibly useful tool. Geosnap has two ways of harmonizing boundaries: areal interpolation and dasymetric interpolation.

With the areal interpolation method, Geosnap uses the area of overlap between polygons of different time periods to create a weighted sum of the overlapped area. The defining factor of this approach compared to dasymetric interpolation is that it **assumes a constant density for each polygon**. In the case of census areas, this would mean that the population is evenly spread throughout the census area. Therefore, this method is best used on small, homogenous polygons (urban areas).



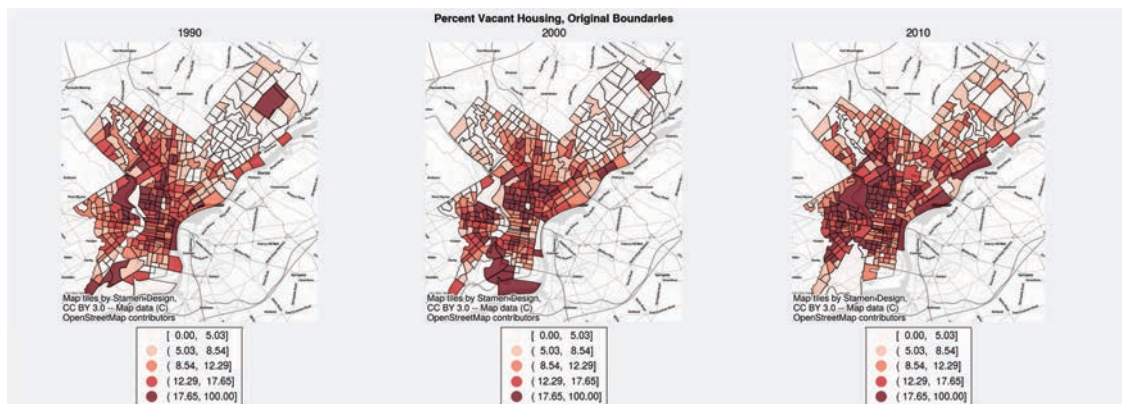
As with many spatial analysis procedures, we must set a reasonable coordinate reference system for the area of focus. In this case, we will use European Petroleum Survey Group (EPSG) 2272.

```
[9]: phl.gdf = phl.gdf.to_crs(2272)
```

To show the Philadelphia census tract boundaries for each time period before harmonization, we will create a time series showing choropleth maps representing a variable for the available census years. The first argument for this command is the name of the column of interest written as a string. `cmap` refers to the map color, which can be assigned to a theme such as “Reds”. `dpi` sets the resolution of the plot, `edgecolor` creates and colors boundary lines for each census tract polygon. `alpha` sets the opacity of the choropleth over the basemap, on a scale from 0-1. `figsize` sets the scale of the plot using inch units (w,h).

```
[10]: phl.plot_timeseries(
        "p_vacant_housing_units", title="Percent Vacant Housing, Original_
        ↪Boundaries",
        cmap='Reds',
        dpi=200,
        edgecolor = 'black',
        alpha = 0.7,
        figsize=(14,5)
    )
```

```
[10]: SubplotsContainer([CartesianAxesSubplot(0.0537579,0.260667;0.225818x0.690944),
    CartesianAxesSubplot(0.387085,0.260667;0.225829x0.690944),
    CartesianAxesSubplot(0.722026,0.260667;0.222615x0.690944)])
```



The output displays three plots showing percent vacant housing per census tract in Philadelphia using each year’s respective census demarcation. You can see that from 1990 to 2010, as the population grew, census tracts were subdivided in some cases, creating new boundaries and therefore hard-to-track demographic changes.

In order to show change more accurately, we will harmonize the data for all three years based on the 2010 boundaries. It is important to note the intensive (the magnitude of which is independent

of the size of the system aka mean, median, temperature, other ratios, etc) and extensive (add to or subtract from the existing system aka population, counts, etc) variables. `weights_method` assigns the interpolation type, and `target_year` sets the desired set of polygons for comparison across time.

```
[11]: phl_harm10 = phl.harmonize(  
        intensive_variables=["p_vacant_housing_units"],  
        extensive_variables=["n_total_pop"],  
        weights_method="area",  
        target_year=2010,  
    )
```

Converting 2 time periods: 0%| | 0/2 [00:00<?, ?it/s]

Harmonizing 1990

C:\Miniconda3\envs\gus5031\lib\site-packages\tobler\util\util.py:28:

UserWarning: nan values in variable: p\_vacant\_housing\_units, replacing with 0  
warn(f"nan values in variable: {column}, replacing with 0")

Harmonizing 2000

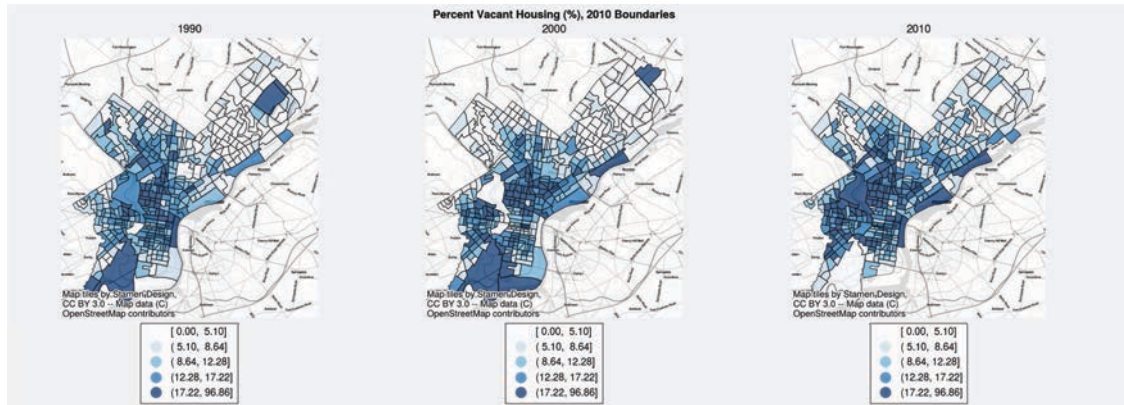
C:\Miniconda3\envs\gus5031\lib\site-packages\tobler\util\util.py:28:

UserWarning: nan values in variable: p\_vacant\_housing\_units, replacing with 0  
warn(f"nan values in variable: {column}, replacing with 0")

Plot the resulting maps:

```
[12]: phl_harm10.plot_timeseries(  
        "p_vacant_housing_units", title="Percent Vacant Housing (%), 2010_  
        ↪Boundaries",  
        dpi=200,  
        edgecolor = 'black',  
        alpha = 0.7,  
        figsize=(14,5)  
    )
```

```
[12]: SubplotsContainer([CartesianAxesSubplot(0.053474,0.260667;0.226385x0.690944),  
    CartesianAxesSubplot(0.386807,0.260667;0.226385x0.690944),  
    CartesianAxesSubplot(0.721674,0.260667;0.223319x0.690944)])
```



## 1 Part II: Modelling Neighborhood & Neighborhood Types

Though the importance of the neighborhood and its centrality to urban life and spatial analysis is nearly indisputable, defining “neighborhood” is tricky for many reasons. They’re wildly complex, different sizes and shapes, and they contain disparate and incomparable objects, with interactions and networks between them that increase exponentially (Claude Fisher). What’s more, a countless number of parties and stakeholders, often with opposing interests and yielding unequal degrees of power, are involved in the process. In Philadelphia, given the city’s age, size, and demographics, these challenges can be both more formidable and more urgent. A 2019 Philadelphia Department of Public Health report, for example, outlined the difficulty the department and its institutional partners had in addressing the health needs of city residents because of the different ways other entities defined and characterized neighborhoods; just months later covid-19 proved their worst-case scenarios were grave underestimates. Researchers faced with these challenges need better tools with which to analyze and make sense of neighborhood change and dynamics both temporally and spatially, as a recent co-authored article by members of the geosnap team claimed (Knaap, 121). And that’s where geosnap comes in: it provides various tools that aim to clarify and measure neighborhood dynamics. In this section, we’ll explore what it provides to analyze **geodemographics**, or the population characteristics that can be divided into geographical areas, and **regionalization**, or the process of identifying and delineating the neighborhoods using those characteristics. GEOSNAP provides both “classic” and spatial clustering algorithms in order to do both. Below, we’ll create a Community object again and use some of geosnap’s algorithms as methods on it.

In the above portion of the tutorial we used census data to create a Community object; here we’ll use the Longitudinal Tract Database (LTDB). Although the LTDB data is derived from the census, I’m using it because it is already harmonized (though this might not be as significant as it appears), it contains data going back to 1970, and it gives us the chance to use different data.

Let’s start with our imports.

```
import os
import sys
import matplotlib.pyplot as plt
import seaborn as sns
from geosnap import Community, datasets
```

## 2 Modelling Neighborhood Types

The `geosnap.analyze` module has a `cluster` function which takes: 1. a list of columns/variables, 2. a clustering algorithm, and 3. a number of clusters.

It splits the dataset by decade and standardizes each variable. It then “re-pools” the time periods back together and runs an algorithm. In doing so, it shows how the distribution of each variable evolves over time. For this part of my demo, I used variables that could indicate the historical effects of redlining, with the aim of importing redlining data into `geosnap` for comparison (which is possible, but beyond our scope).

Begin by creating the `Community` object ‘`phl`’.

```
phl = Community.from_ltdb(county_fips='42101')
```

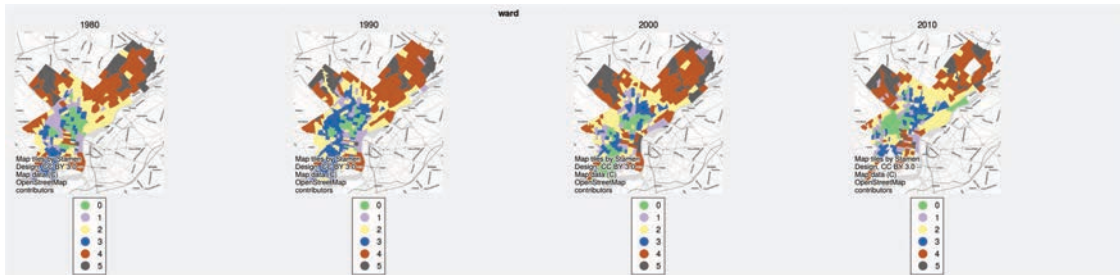
Below, we’ll use the **ward** clustering method, which measures cluster proximity based on the least increase in the difference between the sum of squares in the joint cluster and the combined sum of squares in the given clusters. (Ward’s method is similar to k-means, although it is thought by some to be more accurate regarding clusters of varying physical sizes and clusters “thrown about space very irregularly.”) The `cluster` method returns a `Community` class with cluster labels appended as a new column so we can visualize how the clusters have evolved in space over time.

```
phl = phl.cluster(columns=[
    "median_household_income",
    "p_poverty_rate",
    "p_vacant_housing_units",
    "p_unemployment_rate",
],
    method="ward",
)
```

We’ll plot it with the parameter “`categorical=True`” (as opposed to continuous) and the built-in `map`.

```
phl.plot_timeseries(
    "ward", years=[1980, 1990, 2000, 2010], categorical=True, cmap="Accent"
)
```

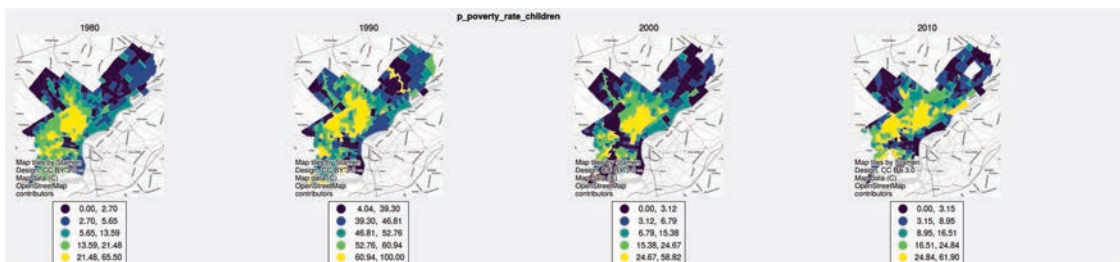
```
SubplotsContainer([CartesianAxesSubplot(0.0280197,0.351445;0.139451x0.591327),
    CartesianAxesSubplot(0.301884,0.351445;0.139451x0.591327),
    CartesianAxesSubplot(0.575748,0.351445;0.139451x0.591327),
    CartesianAxesSubplot(0.851416,0.351445;0.137563x0.591327)])
```



Because cluster labels are appended to the input database we can compute and visualize statistics by cluster, even with variables we didn't use to derive the clusters. Below, I map children's poverty rate (viridis is a built-in map in matplotlib) and then compute the median of the ward clusters using the variable 'p\_recent\_immigrant\_pop'.

```
phl.plot_timeseries(
    "p_poverty_rate_children", years=[1980, 1990, 2000, 2010], cmap="viridis"
)
```

```
SubplotsContainer([CartesianAxesSubplot(0.0280197,0.322186;0.139451x0.618005),
CartesianAxesSubplot(0.301884,0.322186;0.139451x0.618005),
CartesianAxesSubplot(0.575748,0.322186;0.139451x0.618005),
CartesianAxesSubplot(0.851416,0.322186;0.137563x0.618005)])
```



```
phl.gdf.groupby("ward")
```

```
[ "p_recent_immigrant_pop" ].median()[52]: ward
0    0.650289
1    1.882773
2    2.178842
3    1.366304
4    1.816481
5    1.169264
Name: p_recent_immigrant_pop, dtype: float64
```

## 2.1 Clustering Algorithms in GEOSNAP

GEOSNAP provides other clustering algorithms in addition to ward that we can pass to `geosnap.analyze.cluster` to compare. `kmeans`, `affinity_propagation`, `gaussian_mixture`, `spectral` and `hdbscan` are the “classic” ones provided; the spatial algorithms are `azp`, `kmeans_spatial`, `ward_spatial`, `max_p`, `skater`, and `spenc`. (See the `geosnap` api for further info <https://spatialucr.github.io/geosnap/api.html>) We’re not going to use the `hdbscan` because it requires installation of another package, but below I use a **for loop** to compare our data using 3 different algorithms.

```
types = ["ward", "kmeans", "affinity_propagation"]
```

```
for algo_type in types:

    phl = phl.cluster(
        columns=[
            "median_household_income",
            "p_poverty_rate",
            "p_vacant_housing_units",
            "p_unemployment_rate",
        ],
        method=algo_type,
    )
    phl.plot_timeseries(
        column=algo_type,
        years=[1980, 1990, 2000, 2010],
        cmap="Accent",
        categorical=True,
        legend=False,
        alpha=0.6
    )
```

```
C:\Users\tuh42842\AppData\Local\Continuum\anaconda3\envs\gus5031\lib\site-
packages\sklearn\cluster\_kmeans.py:786: FutureWarning: 'precompute_distances'
was deprecated in version 0.23 and will be removed in 1.0 (renaming of 0.25). It
has no effect
```

```
warnings.warn("'precompute_distances' was deprecated in version ")
C:\Users\tuh42842\AppData\Local\Continuum\anaconda3\envs\gus5031\lib\site-
packages\sklearn\cluster\_affinity_propagation.py:148: FutureWarning:
'random_state' has been introduced in 0.23. It will be set to None starting from
1.0 (renaming of 0.25) which means that results will differ at every function
call. Set 'random_state' to None to silence this warning, or to 0 to keep the
behavior of versions <0.23.
warnings.warn(
```





## 2.2 EXERCISES

1. (Easier) Use a different clustering algorithm and/or different variables to compare. One way to do this is to use different variables in the for-loop Try: “gaussian\_mixture” and “spectral.”
2. (More challenging) Since you have already harmonized the census data, go back and model neighborhoods using a regionalization approach. Geosnap will aggregate them into demarcated neighborhoods with internal similarity. To do this, use the **cluster\_spatial** method instead of the classic one we used above. If you are working with larger areas, this might be very helpful to you. Start by creating the new community from the census with the code below:

```
#phl = Community.from_census(county_fips='42101')
```



# 1 Part III

## 1.1 Simulating Neighborhood Sociospatial Dynamics

One of the aspects of geosnap I found exciting is that it allows the user to simulate future neighborhood change. It uses a **spatially-conditioned Markov transition model** which measures change by determining what a certain condition will become based on the distribution of prior transitions between what it was and what it became. Using neighborhood types, geosnap can predict what type a neighborhood is likely to be in the future, provided we know what type a neighborhood was as well as the types of its surrounding neighborhoods.

First, import modules and built in datasets.

```
from geosnap import datasets, Community
import matplotlib.pyplot as plt
```

## 1.2 Cluster Model

As we did above, we start by developing a cluster model of neighborhood types. I chose as variables proven risk factors for severe or terminal cases of covid and/or common characteristics of persons who have had severe or fatal cases of covid using a number of peer-reviewed secondary sources from the past few months. My aim was to get a sense of how so many Philadelphians' vulnerability to and death from covid has accumulated over time as well as to be able to compare actual covid maps with geosnap's simulations. As I mentioned above, the city's Department of Public Health struggled to capture and analyze neighborhood data, and I can't help but think that had we been able to model this better in the past, the city could have targeted populations most at risk for the severe effects of highly communicable infectious diseases. We know this is not the last one we'll experience in our lifetimes; perhaps we can find a way to mitigate its effects before the next one arrives. (Obviously, this is not intended to be "scientific" in any real sense, not least because I had to make decisions about what variables and indicators were in the LTDB that approximated researchers' findings.)

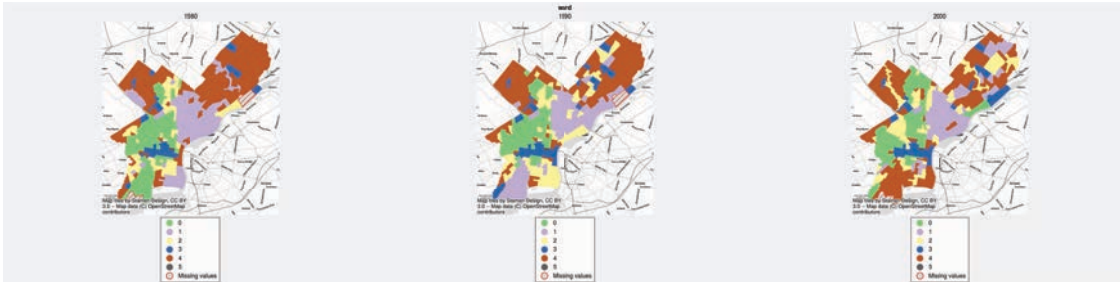
Because we're using the LTDB again, we don't have as many variable options as in the census, but we do have an extra historical decade, which will be important for our simulation below.

```
[ VARS = ['p_black_over_60', 'p_asian_over_60', 'p_housing_units_multiunit_structures', 'p_employed_manufacturing', 'p_disabled']
```

```
phl_clust = Community.from_ltdb(county_fips='42101', years = [1980, 1990, 2000])
```

```
phl_clust = phl_clust.cluster(method='ward', k=6, columns=VARS)
phl_clust.plot_timeseries('ward', categorical=True, figsize=(24,6))
```

```
SubplotsContainer([CartesianAxesSubplot(0.0885011,0.278056;0.156331x0.68162),
CartesianAxesSubplot(0.421834,0.278056;0.156331x0.68162),
CartesianAxesSubplot(0.755168,0.278056;0.156331x0.68162)])
```



## 2 Transition Model ¶

Now that we have the cluster model for the neighborhood types, we can create the Markov transition model which shows which neighborhoods are likely to transition into different neighborhood types. This is crucial for our simulation.

```
phl_clust.plot_transition_matrix('ward', n_rows=2, n_cols=4, figsize=(16,8))
```

## 3 Simulating Neighborhood Change ¶

Now let's simulate the changes. We'll use geosnap's predict method, which takes a cluster model name and simulates it forward by the number of time steps we specify. It returns a new Community. We'll use queen contiguity (which, as we learned last week, defines neighbors as spatial units with a common edge or a common vertices) and 4 time stamps.

```
predicted = phl_clust.simulate('ward', base_year=2000, w_type='queen',
↪time_steps=4)
```

```
C:\Users\tuh42842\AppData\Local\Continuum\anaconda3\envs\gus5031\lib\site-
packages\geopandas\geodataframe.py:1322: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

```
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    super(GeoDataFrame, self).__setitem__(key, value)
C:\Users\tuh42842\AppData\Local\Continuum\anaconda3\envs\gus5031\lib\site-
packages\geopandas\geodataframe.py:1322: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

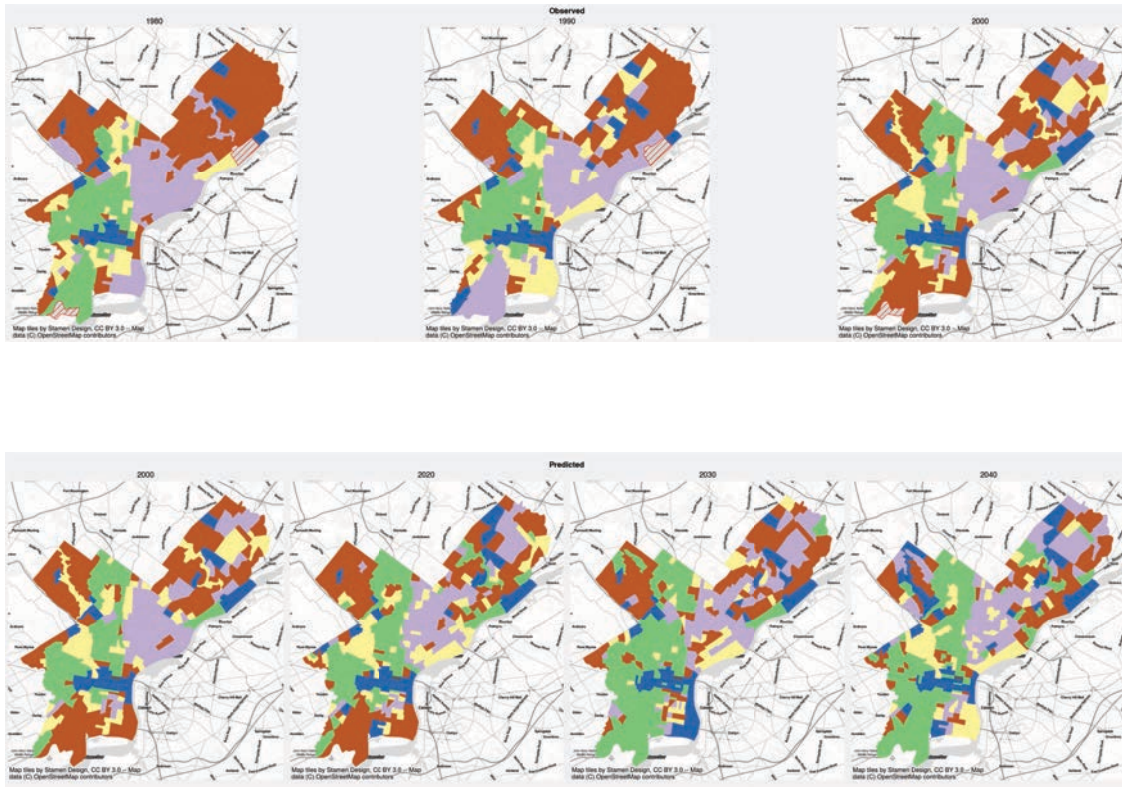
```
See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    super(GeoDataFrame, self).__setitem__(key, value)
C:\Users\tuh42842\AppData\Local\Continuum\anaconda3\envs\gus5031\lib\site-
packages\geopandas\geodataframe.py:1322: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    super(GeoDataFrame, self).__setitem__(key, value)
C:\Users\tuh42842\AppData\Local\Continuum\anaconda3\envs\gus5031\lib\site-
packages\geopandas\geodataframe.py:1322: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
    super(GeoDataFrame, self).__setitem__(key, value)
```

```
phl_clust.plot_timeseries('ward', categorical=True, legend=False,
    ↳ figsize=(20,6), title='Observed')
predicted.gdf = predicted.gdf[predicted.gdf.year!=2010] # the geosnap people
    ↳ say not to use the base year you used above here
predicted.plot_timeseries('ward', categorical=True,
    ↳ legend=False, figsize=(20,6), title='Predicted')
```

```
SubplotsContainer([CartesianAxesSubplot(0.0230417,0.0719981;0.236687x0.859985),
CartesianAxesSubplot(0.268757,0.0719981;0.236687x0.859985),
CartesianAxesSubplot(0.514472,0.0719981;0.236687x0.859985),
CartesianAxesSubplot(0.760188,0.0719981;0.236687x0.859985)])
```



Now, I could conceivably compare the 2020 simulation above with actual covid maps.

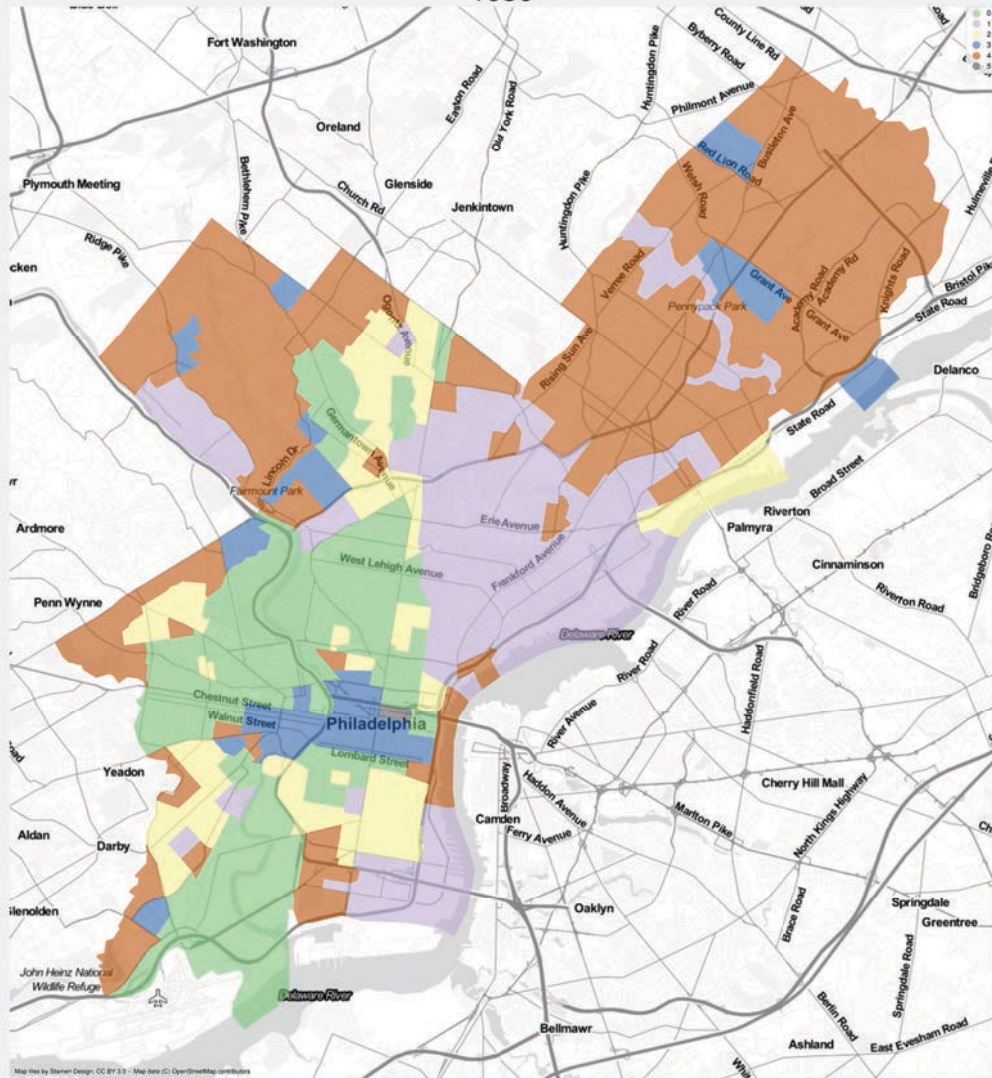
### 3.1 Lastly, let's use geosnap's animate method to put it in a gif.

```
phl_clust.gdf = phl_clust.gdf[phl_clust.gdf.year!= 1990]
combined_sequence = Community.from_geodataframes([phl_clust.gdf, predicted.gdf])
```

In the line below, the filename should be a path to where you will store the gif file. (Note that the execution of this will

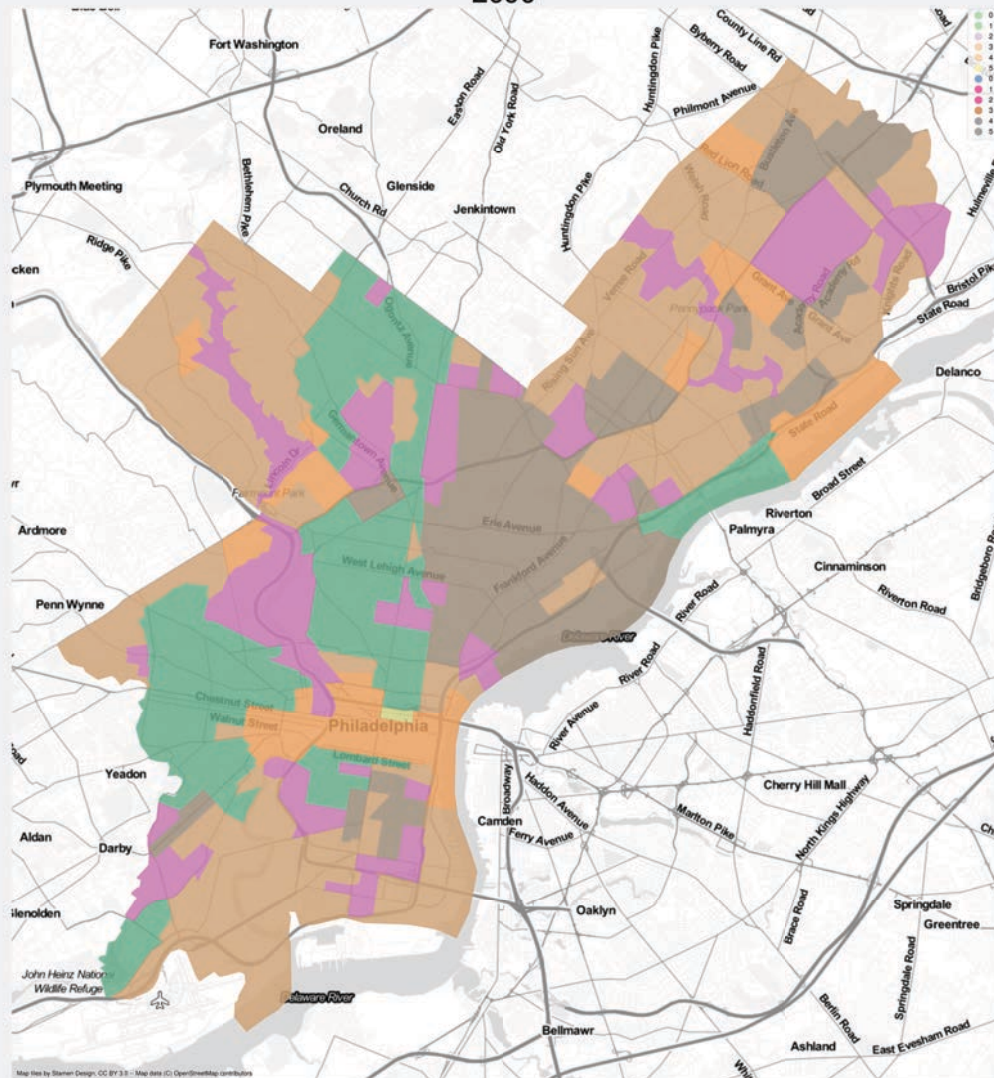
```
combined_sequence.animate_timeseries('ward', categorical = True, filename='C://
↳geosnap-master/examples/images/phl_sim.gif', dpi=200)
```

1980

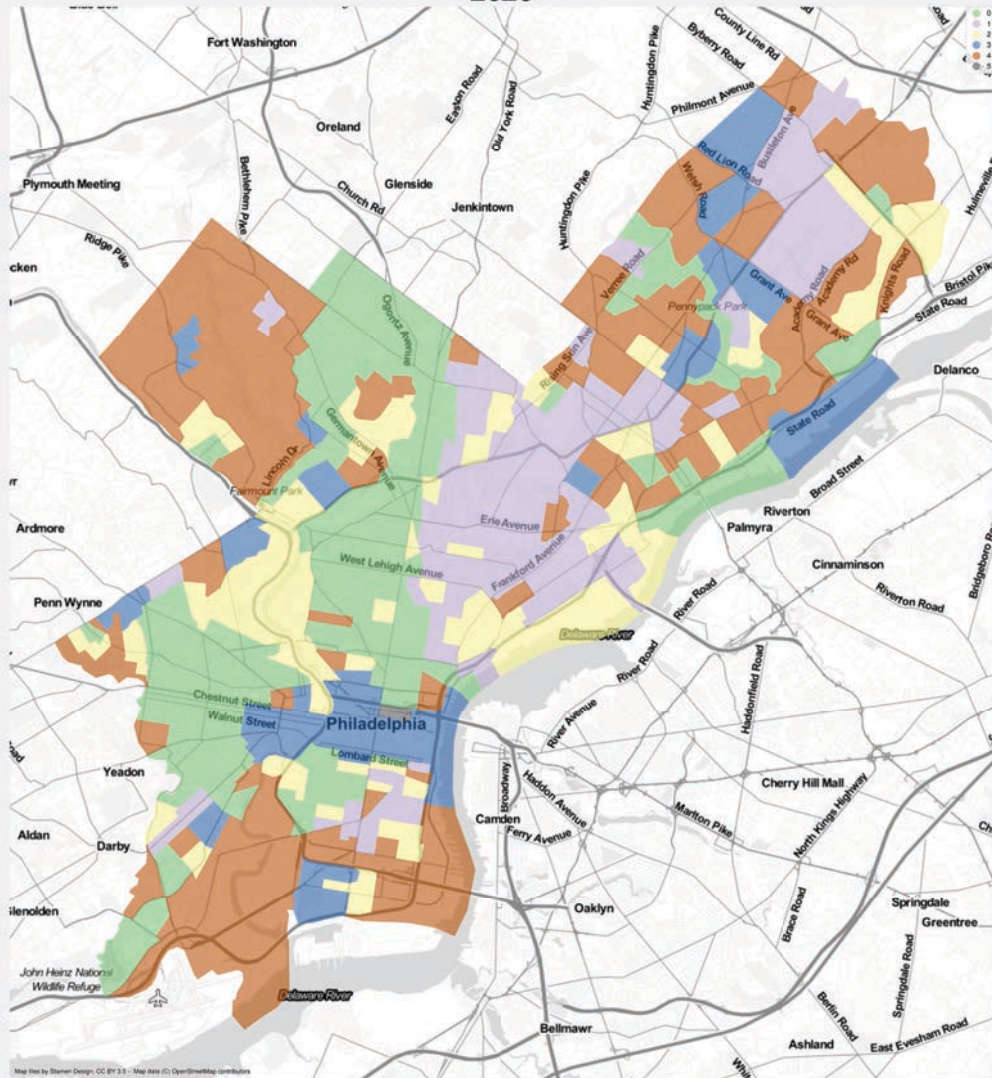




2000

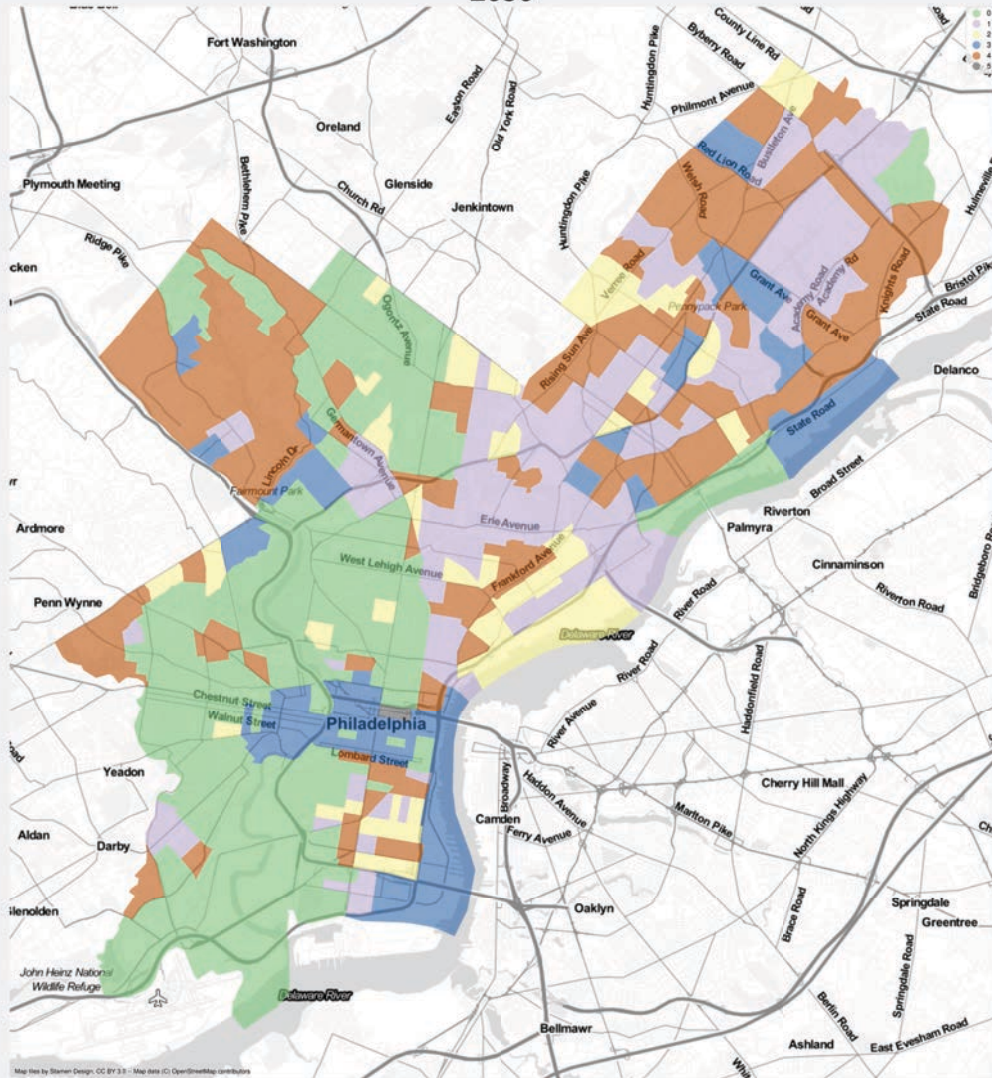


2020



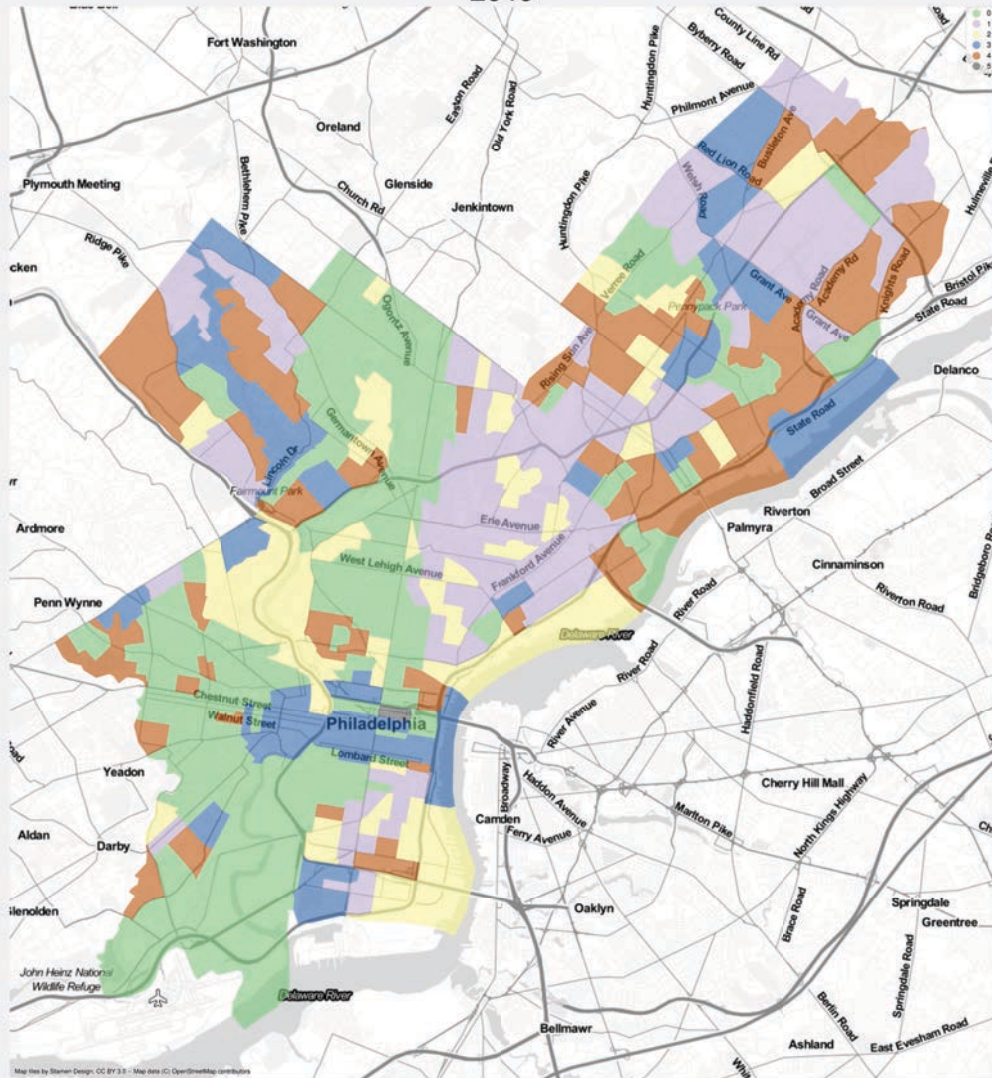


2030





2040





## 1 WORKS CITED

- An, Li, and Stephen Crook. “Spatiotemporal Analysis.” *International Encyclopedia of Geography*, 1–13. American Cancer Society, 2017. <https://doi.org/10.1002/9781118786352.wbieg0635>.
- Carethers, J. M. “Insights into Disparities Observed with COVID-19.” *Journal of Internal Medicine* 289, no. 4 (2021): 463–73. <https://doi.org/10.1111/joim.13199>.
- Fischer, Claude S. “Toward a Subcultural Theory of Urbanism.” *American Journal of Sociology* 80, no. 6 (1975): 1319–41.
- Gao, Ya-dong, Mei Ding, Xiang Dong, Jin-jin Zhang, Ahmet Kursat Azkur, Dilek Azkur, Hui Gan, et al. “Risk Factors for Severe and Critically Ill COVID-19 Patients: A Review.” *Allergy* 76, no. 2 (2021): 428–55. <https://doi.org/10.1111/all.14657>.
- Gerwen, Maaïke, Mathilda Alsen, Christine Little, Joshua Barlow, Eric Genden, Leonard Nymagong, and Douglas Tremblay. “Risk Factors and Outcomes of COVID-19 in New York City; a Retrospective Cohort Study.” *Journal of Medical Virology* 93, no. 2 (February 2021): 907–15. <https://doi.org/10.1002/jmv.26337>.
- Knaap, Elijah, Levi Wolf, Sergio Rey, Wei Kang, and Su Han. “The Dynamics of Urban Neighborhoods: A Survey of Approaches for Modeling Socio-Spatial Structure.” *Journal of GIS & Quantitative Geography* (2019). <https://doi.org/10.31235/osf.io/3frcz>.
- Knaap, Elijah, Wei Kang, Sergio Rey, Levi John Wolf, Renan Xavier Cortes, and Su Han. *geosnap: The Geospatial Neighborhood Analysis Package*. 2019. doi = {10.5281/ZENODO.3526163}.
- Logan, John R., Zengwang Xu, and Brian Stults. “Interpolating US Decennial Census Tract Data from as Early as 1970 to 2010: A Longitudinal Tract Database.” *Professional Geographer* (2012) 66,3: 412–420.
- Mitchell, Andy. *The ESRI Guide to GIS Analysis Volume 2: Spatial Measurements and Statistics*. First edition. Redlands, CA: ESRI, 2005.
- Moore, Kari. “Philadelphia Department of Public Health (PDPH) 500 Cities Project: Linking Local Data for Targeted Action Among Philadelphia’s Most Vulnerable Populations.” Philadelphia Department of Public Health, Urban Health Collaborative, Dornsife School of Public Health, and Drexel University, 2019.
- Rey, Sergio, Su Yeon Han, Wei Kang, Elijah Knaap, and Renan Xavier Cortes. “A Visual Analytics System for Space–Time Dynamics of Regional Income Distributions Utilizing Animated

Flow Maps and Rank-Based Markov Chains.” *Geographical Analysis* 52, no. 4 (2020): 537–57. <https://doi.org/10.1111/gean.12239>.

Silver, Daniel, and Thiago H. Silva. “A Markov Model of Urban Evolution: Neighbourhood Change as a Complex Process.” *PLOS ONE* 16, no. 1 (January 15, 2021). <https://doi.org/10.1371/journal.pone.0245357>.

Soni, Devi. “Introduction to Markov Chains.” *Towards Data Science* (March 5, 2019). Medium.

Steuteville, Robert. “The Once and Future Neighborhood.” *Public Square* (JAN. 29, 2019). Washington, DC: Congress for the New Urbanism.